

Vinamilk (VNM) STOCK PREDICTION

1. Import Necessary Library

```
In [ ]: from vnstock3 import Vnstock
import mplfinance as mpf
import statsmodels.api as sm
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import tensorflow as tf
from tensorflow.keras.layers import Dense, Dropout, Embedding, Bidirectional
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.models import Sequential
import warnings
warnings.filterwarnings('ignore')
import plotly.io as pio
pio.renderers.default = "notebook_connected"
```

2. Exploratory Data Analysis

```
In [ ]: company = Vnstock().stock(symbol='VNM', source='TCBS').company
company.overview()
```

```
Out [ ]:      exchange  industry  company_type  no_shareholders  foreign_percent  outstand
0      HOSE      Thực phẩm và đồ uống      CT      0      0.517
```

```
In [ ]: stock = Vnstock().stock(symbol='VNM', source='VCI')
df = stock.quote.history(start='2019-01-01', end='2024-11-22')
df.info()
df.head()
```

2024-12-04 13:36:35 - vnstock3.common.data.data_explorer - WARNING - Thông tin niêm yết & giao dịch sẽ được truy xuất từ TCBS

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1473 entries, 0 to 1472
Data columns (total 6 columns):
#   Column  Non-Null Count  Dtype
---  -
0   time    1473 non-null    datetime64[ns]
1   open    1473 non-null    float64
2   high    1473 non-null    float64
3   low     1473 non-null    float64
4   close   1473 non-null    float64
5   volume  1473 non-null    int64
dtypes: datetime64[ns](1), float64(4), int64(1)
memory usage: 69.2 KB
```

Out []:

	time	open	high	low	close	volume
0	2019-01-02	78.08	79.69	77.76	79.69	403570
1	2019-01-03	79.69	79.69	78.15	78.98	449730
2	2019-01-04	78.72	80.33	77.76	80.33	498470
3	2019-01-07	82.07	84.19	80.65	84.19	897430
4	2019-01-08	84.44	84.44	82.71	83.54	448350

```
In [ ]: print(df.isnull().sum())

time      0
open      0
high      0
low       0
close     0
volume    0
dtype: int64
```

```
In [ ]: df.describe(include='all')
```

Out []:

	time	open	high	low	close	
count	1473	1473.000000	1473.000000	1473.000000	1473.000000	1
mean	2021-12-14 14:21:15.763747328	73.870754	74.521704	73.213530	73.813327	2
min	2019-01-02 00:00:00	55.810000	57.870000	55.810000	55.810000	2
25%	2020-06-26 00:00:00	66.410000	67.000000	65.810000	66.290000	1
50%	2021-12-10 00:00:00	72.440000	73.200000	71.790000	72.360000	2
75%	2023-06-07 00:00:00	79.980000	80.630000	79.600000	80.050000	3
max	2024-11-22 00:00:00	96.810000	98.070000	95.900000	97.430000	
std	NaN	8.699546	8.705796	8.639016	8.693214	

```
In [ ]: #Filter VNM stock historical data from 01/2024 - 11/2024
#Convert time column
df_analyze = df.copy()
df_analyze['time'] = pd.to_datetime(df_analyze['time'])

#Set time column as index
df_analyze.set_index('time', inplace=True)

#Filter to Plotting chart
filtered_df = df_analyze.loc['2024-01-01':'2024-11-22']

#Plotting the candlestick chart for historical data from 01/01/2024 - 22/
mpf.plot(filtered_df, type='candle', volume=True, style='yahoo', title='V
ylabel='Price', ylabel_lower='Volume', figsize=(14, 7))
```

Vinamilk Stock Price (2024-10-01 to 2024-11-22)



```
In [ ]: # Plot the closing price over time
fig = px.line(x= df_analyze.index, y= df_analyze['close'], title = 'Vinam
fig.update_layout(height=500, width=750, template="plotly_white", font_co
fig.show()
```

Vinamilk Closing Price Over Time



```
In [ ]: from plotly.subplots import make_subplots
# Dimensions for the scatterplot matrix
dimensions = ["open", "high", "low", "close", "volume"]

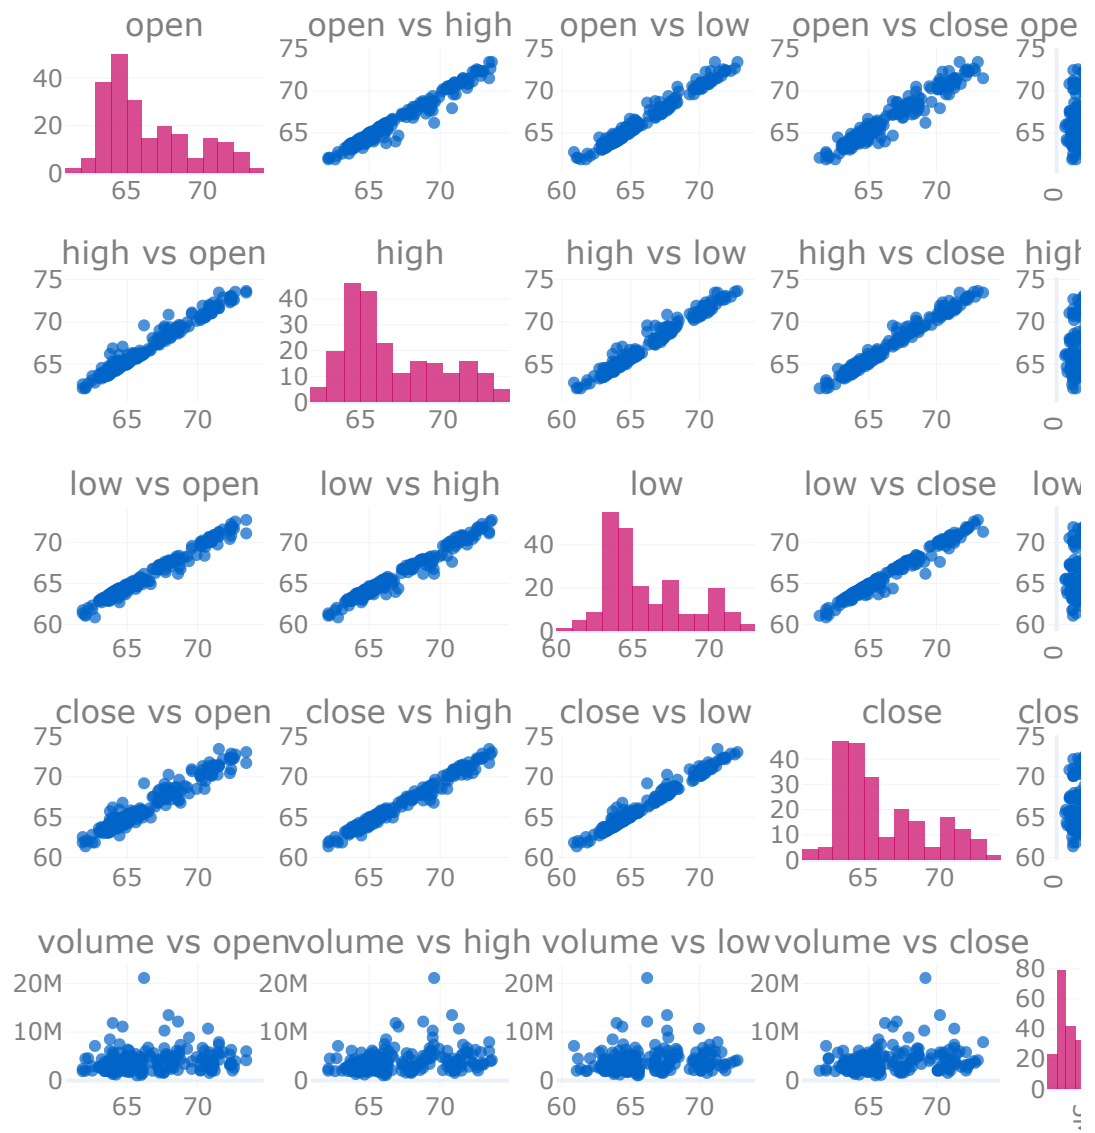
# Create a 5x5 subplot layout
fig = make_subplots(
    rows=len(dimensions), cols=len(dimensions),
    specs=[['type': 'scatter']] * len(dimensions) for _ in range(len(dimensions)),
    subplot_titles=[f"{y} vs {x}" if x != y else f"{x}" for y in dimensions for x in dimensions]
)

# Loop through dimensions to populate the matrix
for i, x in enumerate(dimensions): # Iterate over columns
    for j, y in enumerate(dimensions): # Iterate over rows
        if i != j:
            # Off-diagonal: scatter plot
            fig.add_trace(go.Scatter(x=filtered_df[x], y=filtered_df[y],
                                     marker=dict(size=10, color="blue"),
                                     line=dict(color="blue", width=1)))
        else:
            # Diagonal: histogram
            fig.add_trace(go.Histogram(x=filtered_df[x], marker=dict(color="blue", size=10),
                                     line=dict(color="blue", width=1)))

# Update layout settings
fig.update_layout(height=700, width=750, template="plotly_white", title="Vinamilk Stock Data",
                  font_color="grey", font_size=12, title_font_color="black", title_font_size=14)

# Show the plot
fig.show()
```

Relationship between open, high, low, close, & volume



3. | Modeling (Bidirectional LSTM)

3.1 Preprocessing

```
In [ ]: df_data = df[['time', 'open', 'close']]
df_data.set_index('time', drop=True, inplace=True)
df_data.reset_index(inplace=True)
df_data.info()
df_data.head(10)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1473 entries, 0 to 1472
Data columns (total 3 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   time    1473 non-null    datetime64[ns]
 1   open    1473 non-null    float64
 2   close   1473 non-null    float64
dtypes: datetime64[ns](1), float64(2)
memory usage: 34.7 KB
```

```
Out [ ]:
      time  open  close
0  2019-01-02  78.08  79.69
1  2019-01-03  79.69  78.98
2  2019-01-04  78.72  80.33
3  2019-01-07  82.07  84.19
4  2019-01-08  84.44  83.54
5  2019-01-09  83.54  84.77
6  2019-01-10  84.77  84.64
7  2019-01-11  84.64  86.11
8  2019-01-14  85.47  86.76
9  2019-01-15  86.76  86.69
```

```
In [ ]: # Normalization Data
scaler = MinMaxScaler()
numerical_columns = ['open', 'close']
df_data[numerical_columns] = scaler.fit_transform(df_data[numerical_columns])
```

```
In [ ]: # Split Train - Test data
training_size = round(len(df_data) * 0.80)
train_data = df_data[:training_size]
test_data = df_data[training_size:]

train_data.set_index('time', inplace=True)
test_data.set_index('time', inplace=True)

display(train_data.head())
display(test_data.tail())
display(test_data.shape, train_data.shape)
```

	open	close
time		
2019-01-02	0.543171	0.573763
2019-01-03	0.582439	0.556704
2019-01-04	0.558780	0.589140
2019-01-07	0.640488	0.681884
2019-01-08	0.698293	0.666266

	open	close
time		
2024-11-18	0.192439	0.177559
2024-11-19	0.182683	0.172753
2024-11-20	0.175366	0.187170
2024-11-21	0.192439	0.194378
2024-11-22	0.197317	0.203988

(295, 2)
(1178, 2)

```
In [ ]: def create_sequence(dataset):
    sequences = []
    labels = []
    start_idx = 0
    for stop_idx in range(50, len(dataset)):
        sequences.append(dataset.iloc[start_idx:stop_idx])
        labels.append(dataset.iloc[stop_idx])
        start_idx += 1
    return (np.array(sequences), np.array(labels))

train_seq, train_label = create_sequence(train_data)
test_seq, test_label = create_sequence(test_data)

display(train_seq.shape, test_seq.shape)
```

(1128, 50, 2)
(245, 50, 2)

3.2 Define Bidirectional LSTM

```
In [ ]: def build_BiLSTM_model():
    input = tf.keras.layers.Input(
        shape=(train_seq.shape[1], train_seq.shape[2]), name="input"
    )
    x = tf.keras.layers.Bidirectional(LSTM(512, return_sequences=True))(input)
    x = tf.keras.layers.Bidirectional(LSTM(316, return_sequences=False, dropout=0.5))(x)
    x = tf.keras.layers.Dense(128, activation="relu", name="dense_1")(x)
    output = tf.keras.layers.Dense(2, name="last_dense")(x)

    model = tf.keras.Model(inputs=input, outputs=output)
    model.compile(loss='mean_squared_error', optimizer='adam', metrics=['accuracy'])

    return model

LSTM_model = build_BiLSTM_model()
LSTM_model.summary()
```


Model: "functional"


Layer (type)	Output Shape	
input (InputLayer)	(None, 50, 2)	
bidirectional (Bidirectional)	(None, 50, 1024)	2,1
bidirectional_1 (Bidirectional)	(None, 632)	3,1
dense_1 (Dense)	(None, 128)	
last_dense (Dense)	(None, 2)	


Total params: 5,580,770 (21.29 MB)
Trainable params: 5,580,770 (21.29 MB)
Non-trainable params: 0 (0.00 B)


3.3 Training Data


```
In [ ]: batch_size = 20
        early_stopping_patience = 17
        LSTM_model.fit(train_seq, train_label, epochs=80, validation_data=(test_seq, test_label))
```



Epoch 1/80
36/36  66s 2s/step - loss: 0.0847 - mean_absolute_error: 0.1927 - val_loss: 0.0016 - val_mean_absolute_error: 0.0310


Epoch 2/80
36/36  50s 1s/step - loss: 0.0038 - mean_absolute_error: 0.0472 - val_loss: 8.6478e-04 - val_mean_absolute_error: 0.0230


Epoch 3/80
36/36  45s 1s/step - loss: 0.0024 - mean_absolute_error: 0.0376 - val_loss: 0.0014 - val_mean_absolute_error: 0.0329


Epoch 4/80
36/36  55s 2s/step - loss: 0.0028 - mean_absolute_error: 0.0409 - val_loss: 6.1850e-04 - val_mean_absolute_error: 0.0180


Epoch 5/80
36/36  52s 1s/step - loss: 0.0020 - mean_absolute_error: 0.0337 - val_loss: 5.6969e-04 - val_mean_absolute_error: 0.0177


Epoch 6/80
36/36  54s 2s/step - loss: 0.0021 - mean_absolute_error: 0.0341 - val_loss: 6.1665e-04 - val_mean_absolute_error: 0.0195


Epoch 7/80
36/36  51s 1s/step - loss: 0.0018 - mean_absolute_error: 0.0327 - val_loss: 6.3715e-04 - val_mean_absolute_error: 0.0204


Epoch 8/80
36/36  50s 1s/step - loss: 0.0015 - mean_absolute_error: 0.0292 - val_loss: 5.3048e-04 - val_mean_absolute_error: 0.0168


Epoch 9/80
36/36  53s 1s/step - loss: 0.0016 - mean_absolute_error: 0.0306 - val_loss: 5.0405e-04 - val_mean_absolute_error: 0.0164


Epoch 10/80
36/36  52s 1s/step - loss: 0.0016 - mean_absolute_error: 0.0302 - val_loss: 4.1338e-04 - val_mean_absolute_error: 0.0149


Epoch 11/80
36/36  53s 1s/step - loss: 0.0016 - mean_absolute_error: 0.0309 - val_loss: 4.6872e-04 - val_mean_absolute_error: 0.0158


Epoch 12/80
36/36  48s 1s/step - loss: 0.0017 - mean_absolute_error: 0.0314 - val_loss: 3.6026e-04 - val_mean_absolute_error: 0.0138


Epoch 13/80
36/36  44s 1s/step - loss: 0.0015 - mean_absolute_error: 0.0291 - val_loss: 3.8496e-04 - val_mean_absolute_error: 0.0142


Epoch 14/80
36/36  54s 2s/step - loss: 0.0011 - mean_absolute_error: 0.0249 - val_loss: 3.0202e-04 - val_mean_absolute_error: 0.0124


Epoch 15/80
36/36  42s 1s/step - loss: 0.0011 - mean_absolute_error: 0.0250 - val_loss: 3.7302e-04 - val_mean_absolute_error: 0.0145





















Epoch 16/80
36/36  42s 1s/step - loss: 0.0014 - mean_absolute_error: 0.0285 - val_loss: 3.4260e-04 - val_mean_absolute_error: 0.0139





















Epoch 17/80
36/36  42s 1s/step - loss: 0.0011 - mean_absolute_error: 0.0252 - val_loss: 3.3263e-04 - val_mean_absolute_error: 0.0136





















Epoch 18/80
36/36  42s 1s/step - loss: 0.0011 - mean_absolute_error: 0.0242 - val_loss: 4.0172e-04 - val_mean_absolute_error: 0.0158

Epoch 19/80
36/36  42s 1s/step - loss: 9.6122e-04 - mean_absolute_error: 0.0231 - val_loss: 2.4665e-04 - val_mean_absolute_error: 0.0114

Epoch 20/80
36/36  42s 1s/step - loss: 0.0011 - mean_absolute_error: 0.0247 - val_loss: 2.9030e-04 - val_mean_absolute_error: 0.0121

Epoch 21/80
36/36  42s 1s/step - loss: 0.0011 - mean_absolute_error: 0.0243 - val_loss: 2.4649e-04 - val_mean_absolute_error: 0.0114
Epoch 22/80
36/36  43s 1s/step - loss: 8.9314e-04 - mean_absolute_error: 0.0221 - val_loss: 2.3326e-04 - val_mean_absolute_error: 0.0111
Epoch 23/80
36/36  43s 1s/step - loss: 8.5615e-04 - mean_absolute_error: 0.0215 - val_loss: 3.3937e-04 - val_mean_absolute_error: 0.0140
Epoch 24/80
36/36  43s 1s/step - loss: 0.0011 - mean_absolute_error: 0.0243 - val_loss: 2.1832e-04 - val_mean_absolute_error: 0.0102
Epoch 25/80
36/36  43s 1s/step - loss: 8.6919e-04 - mean_absolute_error: 0.0218 - val_loss: 2.2005e-04 - val_mean_absolute_error: 0.0102
Epoch 26/80
36/36  148s 4s/step - loss: 8.0523e-04 - mean_absolute_error: 0.0208 - val_loss: 5.1243e-04 - val_mean_absolute_error: 0.0198
Epoch 27/80
36/36  42s 1s/step - loss: 8.9428e-04 - mean_absolute_error: 0.0225 - val_loss: 2.2217e-04 - val_mean_absolute_error: 0.0101
Epoch 28/80
36/36  42s 1s/step - loss: 7.8200e-04 - mean_absolute_error: 0.0198 - val_loss: 2.6589e-04 - val_mean_absolute_error: 0.0115
Epoch 29/80
36/36  47s 1s/step - loss: 8.5575e-04 - mean_absolute_error: 0.0213 - val_loss: 4.0100e-04 - val_mean_absolute_error: 0.0163
Epoch 30/80
36/36  43s 1s/step - loss: 8.4289e-04 - mean_absolute_error: 0.0214 - val_loss: 2.2104e-04 - val_mean_absolute_error: 0.0102
Epoch 31/80
36/36  42s 1s/step - loss: 7.8386e-04 - mean_absolute_error: 0.0207 - val_loss: 2.1134e-04 - val_mean_absolute_error: 0.0105
Epoch 32/80
36/36  42s 1s/step - loss: 7.8216e-04 - mean_absolute_error: 0.0204 - val_loss: 3.9510e-04 - val_mean_absolute_error: 0.0164
Epoch 33/80
36/36  42s 1s/step - loss: 7.0554e-04 - mean_absolute_error: 0.0198 - val_loss: 2.0725e-04 - val_mean_absolute_error: 0.0099
Epoch 34/80
36/36  42s 1s/step - loss: 7.0309e-04 - mean_absolute_error: 0.0191 - val_loss: 2.2576e-04 - val_mean_absolute_error: 0.0104
Epoch 35/80
36/36  45s 1s/step - loss: 8.3370e-04 - mean_absolute_error: 0.0211 - val_loss: 2.2169e-04 - val_mean_absolute_error: 0.0105
Epoch 36/80
36/36  41s 1s/step - loss: 7.9471e-04 - mean_absolute_error: 0.0195 - val_loss: 1.9735e-04 - val_mean_absolute_error: 0.0096
Epoch 37/80
36/36  43s 1s/step - loss: 7.9491e-04 - mean_absolute_error: 0.0209 - val_loss: 1.9928e-04 - val_mean_absolute_error: 0.0097
Epoch 38/80
36/36  43s 1s/step - loss: 6.6708e-04 - mean_absolute_error: 0.0188 - val_loss: 2.9345e-04 - val_mean_absolute_error: 0.0134
Epoch 39/80
36/36  44s 1s/step - loss: 7.6442e-04 - mean_absolute_error: 0.0203 - val_loss: 2.0824e-04 - val_mean_absolute_error: 0.0103
Epoch 40/80
36/36  58s 2s/step - loss: 6.0804e-04 - mean_absolute_error: 0.0181 - val_loss: 1.8088e-04 - val_mean_absolute_error: 0.0087

Epoch 41/80
36/36  43s 1s/step - loss: 6.5093e-04 - mean_absolute_error: 0.0185 - val_loss: 2.4430e-04 - val_mean_absolute_error: 0.0114
Epoch 42/80
36/36  43s 1s/step - loss: 9.2134e-04 - mean_absolute_error: 0.0224 - val_loss: 2.6503e-04 - val_mean_absolute_error: 0.0129
Epoch 43/80
36/36  43s 1s/step - loss: 7.1785e-04 - mean_absolute_error: 0.0196 - val_loss: 2.7791e-04 - val_mean_absolute_error: 0.0129
Epoch 44/80
36/36  43s 1s/step - loss: 7.0877e-04 - mean_absolute_error: 0.0199 - val_loss: 2.7300e-04 - val_mean_absolute_error: 0.0119
Epoch 45/80
36/36  43s 1s/step - loss: 7.6886e-04 - mean_absolute_error: 0.0198 - val_loss: 2.3211e-04 - val_mean_absolute_error: 0.0104
Epoch 46/80
36/36  44s 1s/step - loss: 6.9791e-04 - mean_absolute_error: 0.0186 - val_loss: 2.3554e-04 - val_mean_absolute_error: 0.0110
Epoch 47/80
36/36  43s 1s/step - loss: 6.6093e-04 - mean_absolute_error: 0.0188 - val_loss: 2.3564e-04 - val_mean_absolute_error: 0.0120
Epoch 48/80
36/36  43s 1s/step - loss: 7.8752e-04 - mean_absolute_error: 0.0211 - val_loss: 1.8731e-04 - val_mean_absolute_error: 0.0090
Epoch 49/80
36/36  43s 1s/step - loss: 6.2780e-04 - mean_absolute_error: 0.0179 - val_loss: 3.9285e-04 - val_mean_absolute_error: 0.0164
Epoch 50/80
36/36  43s 1s/step - loss: 6.0595e-04 - mean_absolute_error: 0.0180 - val_loss: 1.8783e-04 - val_mean_absolute_error: 0.0090
Epoch 51/80
36/36  43s 1s/step - loss: 7.0762e-04 - mean_absolute_error: 0.0193 - val_loss: 2.0312e-04 - val_mean_absolute_error: 0.0092
Epoch 52/80
36/36  46s 1s/step - loss: 8.2688e-04 - mean_absolute_error: 0.0218 - val_loss: 2.1467e-04 - val_mean_absolute_error: 0.0109
Epoch 53/80
36/36  52s 1s/step - loss: 8.4950e-04 - mean_absolute_error: 0.0215 - val_loss: 2.1890e-04 - val_mean_absolute_error: 0.0103
Epoch 54/80
36/36  42s 1s/step - loss: 6.5526e-04 - mean_absolute_error: 0.0188 - val_loss: 1.9742e-04 - val_mean_absolute_error: 0.0091
Epoch 55/80
36/36  43s 1s/step - loss: 6.6678e-04 - mean_absolute_error: 0.0180 - val_loss: 2.5499e-04 - val_mean_absolute_error: 0.0121
Epoch 56/80
36/36  43s 1s/step - loss: 7.2529e-04 - mean_absolute_error: 0.0197 - val_loss: 2.7022e-04 - val_mean_absolute_error: 0.0123
Epoch 57/80
36/36  43s 1s/step - loss: 7.3742e-04 - mean_absolute_error: 0.0198 - val_loss: 2.1971e-04 - val_mean_absolute_error: 0.0100
Epoch 58/80
36/36  43s 1s/step - loss: 6.2083e-04 - mean_absolute_error: 0.0179 - val_loss: 2.2117e-04 - val_mean_absolute_error: 0.0102
Epoch 59/80
36/36  43s 1s/step - loss: 5.9521e-04 - mean_absolute_error: 0.0177 - val_loss: 2.3059e-04 - val_mean_absolute_error: 0.0104
Epoch 60/80
36/36  83s 1s/step - loss: 7.1195e-04 - mean_absolute_error: 0.0194 - val_loss: 3.1769e-04 - val_mean_absolute_error: 0.0145

Epoch 61/80
36/36  47s 1s/step - loss: 6.4220e-04 - mean_absolute_error: 0.0186 - val_loss: 2.5760e-04 - val_mean_absolute_error: 0.0121
Epoch 62/80
36/36  44s 1s/step - loss: 6.8278e-04 - mean_absolute_error: 0.0187 - val_loss: 2.6474e-04 - val_mean_absolute_error: 0.0118
Epoch 63/80
36/36  44s 1s/step - loss: 8.4559e-04 - mean_absolute_error: 0.0217 - val_loss: 1.7884e-04 - val_mean_absolute_error: 0.0084
Epoch 64/80
36/36  43s 1s/step - loss: 5.3990e-04 - mean_absolute_error: 0.0163 - val_loss: 2.1717e-04 - val_mean_absolute_error: 0.0101
Epoch 65/80
36/36  43s 1s/step - loss: 6.4823e-04 - mean_absolute_error: 0.0181 - val_loss: 2.9531e-04 - val_mean_absolute_error: 0.0135
Epoch 66/80
36/36  1149s 33s/step - loss: 6.6917e-04 - mean_absolute_error: 0.0190 - val_loss: 2.0730e-04 - val_mean_absolute_error: 0.0099
Epoch 67/80
36/36  47s 1s/step - loss: 6.6803e-04 - mean_absolute_error: 0.0185 - val_loss: 2.0730e-04 - val_mean_absolute_error: 0.0092
Epoch 68/80
36/36  43s 1s/step - loss: 6.8554e-04 - mean_absolute_error: 0.0181 - val_loss: 1.9693e-04 - val_mean_absolute_error: 0.0098
Epoch 69/80
36/36  43s 1s/step - loss: 6.3015e-04 - mean_absolute_error: 0.0179 - val_loss: 2.1984e-04 - val_mean_absolute_error: 0.0101
Epoch 70/80
36/36  43s 1s/step - loss: 6.2842e-04 - mean_absolute_error: 0.0185 - val_loss: 4.0819e-04 - val_mean_absolute_error: 0.0167
Epoch 71/80
36/36  43s 1s/step - loss: 6.6370e-04 - mean_absolute_error: 0.0189 - val_loss: 1.9972e-04 - val_mean_absolute_error: 0.0096
Epoch 72/80
36/36  43s 1s/step - loss: 7.8786e-04 - mean_absolute_error: 0.0205 - val_loss: 4.0659e-04 - val_mean_absolute_error: 0.0169
Epoch 73/80
36/36  44s 1s/step - loss: 7.1966e-04 - mean_absolute_error: 0.0194 - val_loss: 3.5757e-04 - val_mean_absolute_error: 0.0147
Epoch 74/80
36/36  43s 1s/step - loss: 7.4310e-04 - mean_absolute_error: 0.0201 - val_loss: 1.8155e-04 - val_mean_absolute_error: 0.0085
Epoch 75/80
36/36  43s 1s/step - loss: 5.6573e-04 - mean_absolute_error: 0.0166 - val_loss: 1.8583e-04 - val_mean_absolute_error: 0.0088
Epoch 76/80
36/36  43s 1s/step - loss: 5.7012e-04 - mean_absolute_error: 0.0173 - val_loss: 2.0073e-04 - val_mean_absolute_error: 0.0097
Epoch 77/80
36/36  44s 1s/step - loss: 6.1334e-04 - mean_absolute_error: 0.0173 - val_loss: 2.6859e-04 - val_mean_absolute_error: 0.0126
Epoch 78/80
36/36  47s 1s/step - loss: 5.9799e-04 - mean_absolute_error: 0.0175 - val_loss: 2.3385e-04 - val_mean_absolute_error: 0.0110
Epoch 79/80
36/36  43s 1s/step - loss: 5.7978e-04 - mean_absolute_error: 0.0170 - val_loss: 2.2740e-04 - val_mean_absolute_error: 0.0112
Epoch 80/80
36/36  44s 1s/step - loss: 6.0592e-04 - mean_absolute_error: 0.0179 - val_loss: 1.7903e-04 - val_mean_absolute_error: 0.0086

Out []: <keras.src.callbacks.history.History at 0x30f11fcd0>

3.4 Evaluate Model

```
In [ ]: test_predicted = LSTM_model.predict(test_seq)
test_inverse_predicted = scaler.inverse_transform(test_predicted)

test_inverse_predicted[:5]
```

8/8 ————— 5s 360ms/step

```
Out [ ]: array([[64.873665, 64.69259 ],
                [64.00423 , 63.799187],
                [64.67144 , 64.51376 ],
                [65.164474, 65.034805],
                [64.67219 , 64.50622 ]], dtype=float32)
```

```
In [ ]: df_slice_data = pd.concat([df_data .iloc[-245:].copy(),pd.DataFrame(test_
df_slice_data[['open','close']] = scaler.inverse_transform(df_slice_data[
display(df_slice_data.head())
df_slice_data.info()
```

	time	open	close	open_predicted	close_predicted
1228	2023-11-30	64.77	63.83	64.873665	64.692589
1229	2023-12-01	64.39	64.77	64.004227	63.799187
1230	2023-12-04	64.87	65.15	64.671440	64.513763
1231	2023-12-05	65.15	64.49	65.164474	65.034805
1232	2023-12-06	64.49	64.87	64.672188	64.506218

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 245 entries, 1228 to 1472

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	time	245 non-null	datetime64[ns]
1	open	245 non-null	float64
2	close	245 non-null	float64
3	open_predicted	245 non-null	float32
4	close_predicted	245 non-null	float32

dtypes: datetime64[ns](1), float32(2), float64(2)

memory usage: 7.8 KB

```
In [ ]: df_slice_data['time'] = pd.to_datetime(df_slice_data['time'])

#Set time column as index
df_slice_data.set_index('time', inplace=True)
df_slice_data
```

Out []:

	open	close	open_predicted	close_predicted
time				
2023-11-30	64.77	63.83	64.873665	64.692589
2023-12-01	64.39	64.77	64.004227	63.799187
2023-12-04	64.87	65.15	64.671440	64.513763
2023-12-05	65.15	64.49	65.164474	65.034805
2023-12-06	64.49	64.87	64.672188	64.506218
...
2024-11-18	63.70	63.20	63.959431	63.757103
2024-11-19	63.30	63.00	63.396782	63.193104
2024-11-20	63.00	63.60	63.143879	62.945084
2024-11-21	63.70	63.90	63.707897	63.527714
2024-11-22	63.90	64.30	63.999714	63.835651


245 rows x 4 columns

In []:

```
# Plot Actual vs Predicted for open price
fig = go.Figure(data=go.Scatter(x= df_slice_data.index, y= df_slice_data[
fig.add_trace(go.Scatter(x=df_slice_data.index, y=df_slice_data['open_pre

fig.update_xaxes(title_text ="Date")
fig.update_yaxes(title_text ="Stock Price")
fig.update_layout(height=500, width=750, template="plotly_white", title =
    font_color="grey", font_size =12,
    title_font_color="black", title_font_size =24)
fig.show()
```


Actual vs Predicted for open price



```
In [ ]: # Plot Actual vs Predicted for close price
fig = go.Figure(data=go.Scatter(x= df_slice_data.index, y= df_slice_data['close_price']))
fig.add_trace(go.Scatter(x=df_slice_data.index, y=df_slice_data['predicted_close_price']))

fig.update_xaxes(title_text = "Date")
fig.update_yaxes(title_text = "Stock Price")
fig.update_layout(height=500, width=750, template="plotly_white", title = "Actual vs Predicted for close price",
                    font_color="grey", font_size = 12,
                    title_font_color="black", title_font_size = 24,
                    legend_title=None, legend=dict(orientation='h', yanchor='bottom',
fig.show()
```

Actual vs Predicted for close price



4. Further Prediction (using Bidirectional LSTM)

```
In [ ]: # Convert NumPy array to DataFrame
df_data = pd.DataFrame(df_data, columns=['open', 'close'])

# Keep rows from index 0 to 1472
df_data = df_data.iloc[:1473]

# Reset index
df_data.reset_index(drop=True, inplace=True)
# Create a date range for the next 180 business days
last_date = df['time'].iloc[-1]
business_days = pd.date_range(start=last_date + pd.Timedelta(days=1), per

predictions = []

# Loop through the business days
for next_date in business_days:
    # Get the last 50 rows of 'open' and 'close' prices for prediction
    last_50_data = df_data[-50:] # Get the last 50 normalized rows

    # Reshape the data for LSTM input
    x_next = np.array([last_50_data]) # Shape will be (1, 50, 2)
    x_next = np.reshape(x_next, (x_next.shape[0], x_next.shape[1], 2)) #

    # Make predictions using the trained LSTM model
    y_next_predict = LSTM_model.predict(x_next)
```



```
# Inverse transform the predicted values to get actual prices
y_next_predict_inverse = scaler.inverse_transform(y_next_predict)

# Prepare the DataFrame for the next day's prediction
df_next = pd.DataFrame({
    'time': [next_date],
    'open_predicted': [y_next_predict_inverse[0][0]],
    'close_predicted': [y_next_predict_inverse[0][1]]
})

# Append the prediction to the predictions list
predictions.append(df_next)

# Normalize the predicted values before appending to df_data
new_row_normalized = scaler.transform(np.array([[y_next_predict_inverse[0][0], y_next_predict_inverse[0][1]]]))

# Append the new normalized predicted row to the normalized DataFrame
new_row = pd.DataFrame(new_row_normalized, columns=['open', 'close'])
df_data = pd.concat([df_data, new_row], ignore_index=True) # Append

# Combine all predictions into a single DataFrame
df_predictions = pd.concat(predictions, ignore_index=True)

df_slice_col=df_slice_data.reset_index()
final_df = pd.concat([df_slice_col, df_predictions], ignore_index=True)
print(final_df)
```

1/1		0s	57ms/step
1/1		0s	57ms/step
1/1		0s	48ms/step
1/1		0s	50ms/step
1/1		0s	45ms/step
1/1		0s	48ms/step
1/1		0s	49ms/step
1/1		0s	46ms/step
1/1		0s	47ms/step
1/1		0s	49ms/step
1/1		0s	48ms/step
1/1		0s	63ms/step
1/1		0s	55ms/step
1/1		0s	48ms/step
1/1		0s	47ms/step
1/1		0s	44ms/step
1/1		0s	47ms/step
1/1		0s	46ms/step
1/1		0s	45ms/step
1/1		0s	44ms/step
1/1		0s	46ms/step
1/1		0s	58ms/step
1/1		0s	51ms/step
1/1		0s	51ms/step
1/1		0s	48ms/step
1/1		0s	52ms/step
1/1		0s	52ms/step
1/1		0s	50ms/step
1/1		0s	51ms/step
1/1		0s	57ms/step
1/1		0s	65ms/step
1/1		0s	173ms/step
1/1		0s	271ms/step
1/1		0s	186ms/step
1/1		0s	159ms/step
1/1		0s	105ms/step
1/1		0s	127ms/step
1/1		0s	99ms/step
1/1		0s	85ms/step
1/1		0s	100ms/step
1/1		0s	118ms/step
1/1		0s	85ms/step
1/1		0s	53ms/step
1/1		0s	53ms/step
1/1		0s	52ms/step
1/1		0s	48ms/step
1/1		0s	47ms/step
1/1		0s	47ms/step
1/1		0s	45ms/step
1/1		0s	46ms/step
1/1		0s	47ms/step
1/1		0s	50ms/step
1/1		0s	50ms/step
1/1		0s	48ms/step
1/1		0s	49ms/step
1/1		0s	49ms/step
1/1		0s	49ms/step
1/1		0s	49ms/step
1/1		0s	50ms/step

1/1		0s	52ms/step
1/1		0s	50ms/step
1/1		0s	53ms/step
1/1		0s	52ms/step
1/1		0s	54ms/step
1/1		0s	53ms/step
1/1		0s	53ms/step
1/1		0s	82ms/step
1/1		0s	66ms/step
1/1		0s	69ms/step
1/1		0s	55ms/step
1/1		0s	55ms/step
1/1		0s	55ms/step
1/1		0s	63ms/step
1/1		0s	64ms/step
1/1		0s	63ms/step
1/1		0s	55ms/step
1/1		0s	47ms/step
1/1		0s	47ms/step
1/1		0s	47ms/step
1/1		0s	50ms/step
1/1		0s	48ms/step
1/1		0s	50ms/step
1/1		0s	46ms/step
1/1		0s	49ms/step
1/1		0s	51ms/step
1/1		0s	53ms/step
1/1		0s	52ms/step
1/1		0s	48ms/step
1/1		0s	47ms/step
1/1		0s	50ms/step
1/1		0s	48ms/step
1/1		0s	49ms/step
1/1		0s	50ms/step
1/1		0s	53ms/step
1/1		0s	52ms/step
1/1		0s	53ms/step
1/1		0s	54ms/step
1/1		0s	80ms/step
1/1		0s	63ms/step
1/1		0s	55ms/step
1/1		0s	55ms/step
1/1		0s	56ms/step
1/1		0s	57ms/step
1/1		0s	59ms/step
1/1		0s	55ms/step
1/1		0s	56ms/step
1/1		0s	56ms/step
1/1		0s	58ms/step
1/1		0s	56ms/step
1/1		0s	57ms/step
1/1		0s	55ms/step
1/1		0s	57ms/step
1/1		0s	56ms/step
1/1		0s	59ms/step
1/1		0s	57ms/step
1/1		0s	55ms/step
1/1		0s	55ms/step
1/1		0s	54ms/step
1/1		0s	56ms/step

1/1	0s	60ms/step
1/1	0s	57ms/step
1/1	0s	58ms/step
1/1	0s	58ms/step
1/1	0s	57ms/step
1/1	0s	58ms/step
1/1	0s	57ms/step
1/1	0s	55ms/step
1/1	0s	80ms/step
1/1	0s	66ms/step
1/1	0s	55ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	57ms/step
1/1	0s	58ms/step
1/1	0s	55ms/step
1/1	0s	55ms/step
1/1	0s	57ms/step
1/1	0s	54ms/step
1/1	0s	57ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	56ms/step
1/1	0s	58ms/step
1/1	0s	57ms/step
1/1	0s	67ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	56ms/step
1/1	0s	57ms/step
1/1	0s	59ms/step
1/1	0s	61ms/step
1/1	0s	55ms/step
1/1	0s	71ms/step
1/1	0s	66ms/step
1/1	0s	60ms/step
1/1	0s	56ms/step
1/1	0s	56ms/step
1/1	0s	53ms/step
1/1	0s	53ms/step
1/1	0s	57ms/step
1/1	0s	59ms/step
1/1	0s	54ms/step
1/1	0s	56ms/step
1/1	0s	57ms/step
1/1	0s	56ms/step
1/1	0s	56ms/step
1/1	0s	58ms/step
1/1	0s	55ms/step
1/1	0s	58ms/step
1/1	0s	56ms/step
1/1	0s	57ms/step
1/1	0s	55ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	56ms/step
1/1	0s	54ms/step
1/1	0s	57ms/step

```
1/1 ————— 0s 58ms/step
```

	time	open	close	open_predicted	close_predicted
0	2023-11-30	64.77	63.83	64.873665	64.692589
1	2023-12-01	64.39	64.77	64.004227	63.799187
2	2023-12-04	64.87	65.15	64.671440	64.513763
3	2023-12-05	65.15	64.49	65.164474	65.034805
4	2023-12-06	64.49	64.87	64.672188	64.506218
...
420	2025-07-28	NaN	NaN	66.621880	66.604301
421	2025-07-29	NaN	NaN	66.626602	66.609459
422	2025-07-30	NaN	NaN	66.631470	66.614761
423	2025-07-31	NaN	NaN	66.636467	66.620178
424	2025-08-01	NaN	NaN	66.641571	66.625702

[425 rows x 5 columns]

```
In [ ]: file_name = 'VNM 6 months prediction.xlsx'

# saving the excel
final_df.to_excel(file_name)
```

```
In [ ]: print(df_predictions.head(15))
```

	time	open_predicted	close_predicted
0	2024-11-25	64.396118	64.245888
1	2024-11-26	64.347000	64.185394
2	2024-11-27	64.280724	64.114395
3	2024-11-28	64.221153	64.053589
4	2024-11-29	64.181702	64.013023
5	2024-12-02	64.163872	63.993649
6	2024-12-03	64.160667	63.988586
7	2024-12-04	64.162994	63.989716
8	2024-12-05	64.165321	63.992329
9	2024-12-06	64.162216	63.992214
10	2024-12-09	64.153114	63.988899
11	2024-12-10	64.140511	63.984440
12	2024-12-11	64.126968	63.980957
13	2024-12-12	64.116875	63.981350
14	2024-12-13	64.113014	63.985970

```
In [ ]: # Plotting the predicted open data until 7/2025
fig1 = go.Figure(data=go.Scatter(x= final_df['time'], y= final_df['open']
fig1.add_trace(go.Scatter(x=final_df['time'], y=final_df['open_predicted']
                        marker_color ="blue", name='Predicted Open Price


fig1.update_xaxes(title_text ="Date")
fig1.update_yaxes(title_text ="Stock Price")
fig1.update_layout(height=500,width=800, template="plotly_white", title =
                    font_color="grey", font_size =12,
                    title_font_color="black", title_font_size =24)
fig1.show()

# Plotting the predicted close data until 7/2025
fig2 = go.Figure(data=go.Scatter(x= final_df['time'], y= final_df['close']
fig2.add_trace(go.Scatter(x=final_df['time'], y=final_df['close_predicted']
                        marker_color ="blue", name='Predicted Close Pric

fig2.update_xaxes(title_text ="Date")
fig2.update_yaxes(title_text ="Stock Price")
fig2.update_layout(height=500, width=800, template="plotly_white", title
```

```
font_color="grey", font_size =12,  
title_font_color="black", title_font_size =24,  
legend_title=None, legend=dict(orientation='h', yanchor='bottom',  
fig2.show()
```

Vinamilk Stock Open Price Prediction for the r



Vinamilk Stock Close Price Prediction for the I

